Surface Pattern Recognition using Shape Subspace

Ibrahim Naish
Supervisor: Kazuhiro Fukui

January 22, 2025

Abstract

In this study, we propose a method for surface pattern recognition on three-dimensional floor point cloud data by leveraging Shape Subspace methods and the Difference Subspace (DS) technique. The dataset consists of synthetic three-dimensional floor surfaces represented as point clouds, which are segmented into smaller patches for localized analysis. To generate the shape subspace, the point cloud data is first centered by subtracting the mean to remove translational bias, ensuring that the analysis focuses on intrinsic geometric features. Singular Value Decomposition (SVD) is then applied to extract the dominant geometric components of each patch. The Difference Subspace is subsequently computed by comparing the orthogonal projection matrices of neighboring patches, enabling the detection of subtle surface variations. A series of experimental evaluations were conducted using both four-neighbor and eight-neighbor configurations. The objective of these evaluations was to assess the sensitivity and robustness of the proposed approach. The results of these evaluations demonstrate that the four-neighbor configuration provides higher sensitivity to fine-grained surface variations. Conversely, the eight-neighbor approach captures broader contextual differences with improved stability.

1 Introduction

Surface pattern recognition plays a crucial role in various domains, including computer vision, robotics, and quality inspection. In these fields, precise identification of local geometric variations is essential. Traditional methods, including curvature analysis and spatial filtering, often struggle to capture these subtle variations, particularly in intricate three-dimensional surfaces.

Point clouds have emerged as a widely used representation in three-dimensional data processing. This representation consists of a set of points in a three-dimensional coordinate system, typically defined by (x, y, z) coordinates. These point clouds are commonly generated from three-dimensional scanning techniques, such as LiDAR and structured light scanning, making them valuable for capturing detailed and accurate representations of surfaces and objects. However, their unstructured nature introduces limitations in terms of processing and analysis efficiency when compared to structured representations, such as meshes or voxel grids. Additionally, point clouds are susceptible to noise and outliers, which can stem from environmental factors or sensor inaccuracies [5].

Subspace-based methods have demonstrated efficacy in addressing these challenges by providing a robust mathematical framework for analyzing and comparing geometric structures. The concept of shape subspaces [4], enables the comparison of three-dimensional objects by analyzing the geometric relationships between subspaces using canonical angles[1]. These methods have been widely applied in various computer vision tasks, such as multi-view object recognition and facial recognition [3, 6]. The factorization method, a key technique in analysis based on subspaces, decomposes a measurement matrix into motion and shape components, allowing for the extraction of shape subspaces that are resistant to transformations such as rotation and translation. This method ensures consistent feature representation by factorizing a set of tracked feature points into a shape matrix and a motion matrix, making it well-suited for applications involving dynamic or complex geometries. However, in the context of point cloud data, where the three-dimensional coordinates (x, y, z) are directly available, factorization is not required, allowing for direct subspace computation.

Building on this foundation, Difference Subspace (DS) techniques [2], have been introduced as an extension of the concept of a difference vector to subspaces. The DS framework is predicated on the quantification of geometric differences between subspaces by isolating their distinguishing components, thus offering a more refined approach to surface variation analysis. A distinguishing feature of DS is its ability to perform localized comparisons, a capability that surpasses the global comparison methods traditionally employed. This property renders DS particularly effective for detecting subtle shape variations within neighboring regions.

In this study, the DS framework is employed to detect local surface variations in synthetic three-dimensional floor surfaces represented as point clouds. The proposed approach involves the segmentation of the floor into patches and the subsequent application of Singular Value Decomposition (SVD) to generate subspaces for each patch. The DS is then computed between neighboring patches to quantify local shape differences. To assess the sensitivity and robustness of the proposed approach, we evaluated it using both four-neighbor and eight-neighbor configurations. The former provides increased sensitivity to small variations, while the latter captures broader contextual differences.

In the following sections, we first describe the methodology used to generate shape subspaces and compute Difference Subspaces from point cloud data. We then present experimental results comparing different neighborhood configurations and discuss their effectiveness in capturing surface irregularities. Finally, we conclude with a discussion on the implications of our findings and potential future directions, such as extending the approach to real-world datasets and improving noise resilience.

2 Methodology

This research proposes a framework that employs shape subspace methods combined with the Difference Subspace (DS) technique to analyze local surface patterns on three-dimensional floor point cloud data. The methodology consists of several key stages, including data generation, patch division, subspace computation, and difference subspace analysis.

2.1 Data Generation

In order to facilitate the execution of the experiments, a synthetic three-dimensional floor surface is generated with controllable variations such as bumps and ridges as seen in Figure 1. The floor is represented as a point cloud with (x, y, z) coordinates. The synthetic floor is defined by dimensions of 50×50 where each point is spaced exactly one centimeter apart, resulting in a grid of 10,000 points distributed evenly across the surface.

2.2 Patch Division

The generation of point cloud data is followed by its division into smaller patches, a process that facilitates localized analysis. The division process involves splitting the floor into an n x n grid structure, preserving spatial coherence and allowing for efficient subspace computation. The experimental design encompasses a range of patch configurations, including:

- 3×3, 4×4, 5×5, 6×6, 8×8, and 10×10 **grid sizes.**
- Two neighbor comparison approaches: **four-neighbor and eight-neighbor.**

2.3 Generation of Shape Subspace

Point cloud data consists of a set of discrete three-dimensional points, typically represented by their Cartesian coordinates (x, y, z). These points are often obtained from three-dimensional scanning techniques such as LiDAR, structured light, or stereo vision. They represent the surface of an object or environment without any inherent structure or connectivity. Due to the unstructured nature of point clouds, analyzing their geometric properties requires transforming them into a more structured representation, such as a subspace representation.

The shape subspace is generated for each patch independently to capture local geometric properties of the surface. Each patch consists of a set of three-dimensional points (x, y, z), which are extracted from the original point cloud after segmentation into a uniform grid structure. This segmentation allows localized analysis and ensures spatial coherence across the surface.

To generate the shape subspace, the following steps are performed for each patch.

Given n three-dimensional points $\{x_i, y_i, z_i\}_{i=1}^{n}$, the matrix $\mathbf{V}$ is constructed as:

$$\mathbf{V} = \begin{bmatrix} x_1 & y_1 & z_1 \\ x_2 & y_2 & z_2 \\ \vdots & \vdots & \vdots \\ x_n & y_n & z_n \end{bmatrix}.$$

To ensure translation invariance, the mean position of the points is subtracted from each coordinate:

$$\mathbf{V}_{\text{centered}} = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} & z_1 - \bar{z} \\ x_2 - \bar{x} & y_2 - \bar{y} & z_2 - \bar{z} \\ \vdots & \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} & z_n - \bar{z} \end{bmatrix}.$$

Once the data is centered, Singular Value Decomposition (SVD) is applied to analyze the geometric structure of the points:

$$\mathbf{V}_{\text{centered}} = \mathbf{U}\mathbf{\Sigma}\mathbf{\Phi}^T.$$

The basis of the shape subspace is obtained by selecting the first d columns from the left singular matrix $\mathbf{U}$:

$$\mathbf{S} = \mathbf{U}[:, : d].$$

where d is the desired subspace dimension (e.g., 3 for three-dimensional shapes). This selection captures the primary geometric features of the point cloud.

2.4 Generation of Difference Subspace

Once the shape subspaces for all patches have been generated, the next step involves the computation of the Difference Subspace (DS) between neighboring patches to quantify local surface variations. The DS is computed by comparing the orthogonal projection matrices of two adjacent patches. The orthogonal projection matrix onto a given space is uniquely determined by its corresponding orthonormal basis. The objective is to compute the Difference Space (DS) that captures variations between regions, given a set of three-dimensional points. [4]

1. Compute Orthonormal Bases.

   Let $\mathbf{\Phi} = [\Phi_1, \Phi_2, \ldots, \Phi_{N_\mathcal{P}}]$ be an orthonormal basis for $N_\mathcal{P}$-dimensional subspace $\mathcal{P}$, and let $\mathbf{\Psi} = [\Psi_1, \Psi_2, \ldots, \Psi_{N_\mathcal{Q}}]$ be an orthonormal basis for $N_\mathcal{Q}$-dimensional subspace $\mathcal{Q}$.

2. Define Projection Matrices.

   The orthogonal projection matrices onto the subspaces are defined as:

$$\mathbf{P} = \sum_{i=1}^{N_\mathcal{P}} \Phi_i \Phi_i^T = \mathbf{\Phi}\mathbf{\Phi}^T, \quad \mathbf{Q} = \sum_{i=1}^{N_\mathcal{Q}} \Psi_i \Psi_i^T = \mathbf{\Psi}\mathbf{\Psi}^T$$

3. Compute the Sum of Projection Matrices.

   The sum matrix of the projection matrices is calculated as follows:

$$\mathbf{S} = \mathbf{P} + \mathbf{Q}.$$

4. Perform Eigenvalue Decomposition on $\mathbf{S}$.

$$\mathbf{S}\mathbf{v}_i = \lambda_i \mathbf{v}_i.$$

5. Identify the Difference Subspace.

   The eigenvectors corresponding to eigenvalues less than 1 form the basis of the difference subspace:

$$\mathcal{D} = \text{Span}\{\mathbf{v}_i \mid \lambda_i < 1\}.$$

3 Experiment and Discussion

An analysis of patch sizes ranging from 3×3 to 10×10, employing both the four-neighbor and eight-neighbor approaches, reveals noteworthy trends in performance. Generally, smaller patch sizes (e.g., 3×3 and 4×4) demonstrate higher accuracy and lower standard deviation, signifying their aptitude for detecting fine-grained surface variations. The four-neighbor approach consistently delivers reliable results, especially in cases where local variations are significant. However, as patch size increases, standard deviation rises, indicating a reduction in reliability due to the inclusion of broader surface areas, which can introduce more noise and reduce sensitivity to subtle variations.

The findings of this study suggest that the selection of patch size and neighborhood configuration should be contingent upon the specific application in question. In the context of fine-grained surface detection, the utilization of smaller patches with a four-neighbor approach has been demonstrated to yield enhanced sensitivity and stability. Conversely, for broader pattern recognition, the employment of larger patches with an eight-neighbor configuration has been demonstrated to facilitate a more detailed understanding of surface variations.

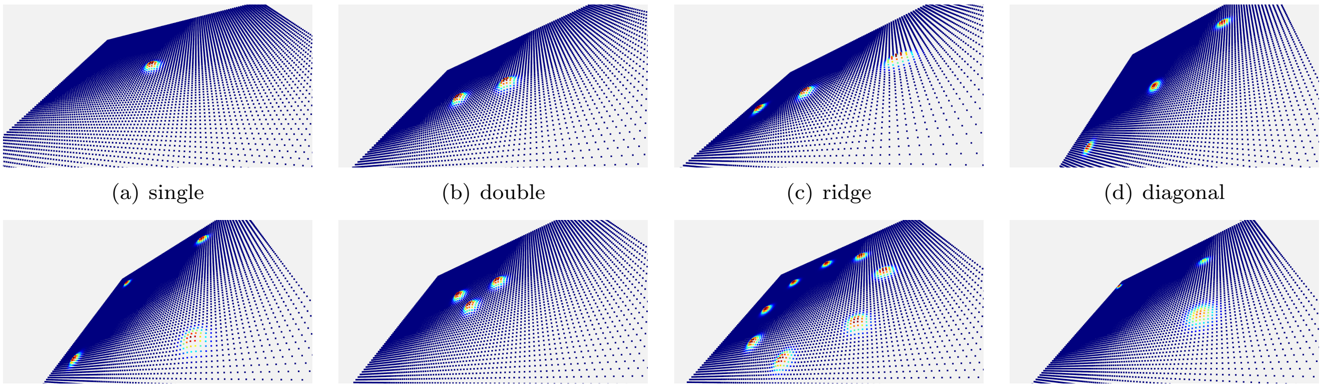

(a) single (b) double (c) ridge (d) diagonal

(e) corners (f) triangle (g) circle (h) random

Figure 1: Generated patterns

4 Conclusion and future work

In this research, we addressed surface pattern recognition on three-dimensional floor point cloud data using Shape Subspace methods and the Difference Subspace (DS) technique. Initially, the floor was segmented into patches to facilitate localized analysis. The centering of each patch was achieved by subtracting the mean, a process that aimed to eliminate positional bias and ensure that the computed subspaces captured intrinsic shape features. Subsequently, the singular value decomposition (SVD) was used to extract the predominant geometric components for each patch.

The DS was computed by comparing the orthogonal projection matrices of neighboring patches to quantify local differences. A comparison of the four-neighbor and eight-neighbor approaches demonstrated that the former provides greater sensitivity to fine-grained variations, while the latter offers improved stability for complex patterns. The effectiveness of these methods in detecting surface irregularities was validated through scatter plots and heat maps.

Future work will focus on applying the method to real-world datasets, enhancing robustness against noise, and optimizing computational performance for larger datasets. Adaptive patch sizing and configuration strategies will also be explored. Additionally, we aim to extend the method to practical applications such as road surface analysis and structural health monitoring to further demonstrate its utility.

References

[1] Kazuhiro Fukui. "Subspace Methods". In: Jan. 2020, pp. 1–5. ISBN: 978-3-030-03243-2. DOI: 10.1007/978-3-030-03243-2_708-1.

[2] Kazuhiro Fukui and Atsuto Maki. "Difference Subspace and Its Generalization for Subspace-Based Methods". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 37.11 (2015), pp. 2164–2177. DOI: 10.1109/TPAMI.2015.2408358.

[3] Kazuhiro Fukui and Osamu Yamaguchi. "Face Recognition Using Multi-viewpoint Patterns for Robot Vision". In: *Robotics Research. The Eleventh International Symposium*. Ed. by Paolo Dario and Raja Chatila. Berlin, Heidelberg: Springer Berlin Heidelberg, 2005, pp. 192–201. ISBN: 978-3-540-31508-7.

[4] Yosuke Igarashi and Kazuhiro Fukui. "3D Object Recognition Based on Canonical Angles between Shape Subspaces". In: *Computer Vision – ACCV 2010*. Ed. by Ron Kimmel, Reinhard Klette, and Akihiro Sugimoto. Berlin, Heidelberg: Springer Berlin Heidelberg, 2011, pp. 580–591. ISBN: 978-3-642-19282-1.

[5] "Point Cloud Processing and Compression". In: *IEEE Transactions on Circuits and Systems for Video Technology* 30.10 (2020), pp. 3856–3856. DOI: 10.1109/TCSVT.2020.3019906.

[6] Takao Yoshinuma, Hideitsu Hino, and Kazuhiro Fukui. "Personal Authentication Based on 3D Configuration of Micro-feature Points on Facial Surface". In: Feb. 2016, pp. 433–446. ISBN: 978-3-319-29450-6. DOI: 10.1007/978-3-319-29451-3_35.